

CHAUFFEUR: Benchmark Suite for Design and End-to-End Analysis of Self-Driving Vehicles on Embedded Systems

CS637A – Course Project

Course Instructor: Dr. Indranil Saha

Group 6

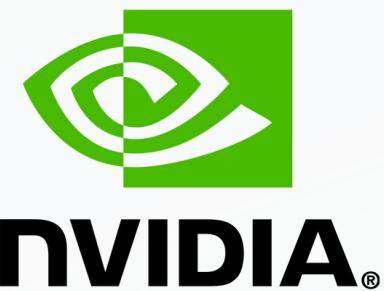
Arnav Pandey (200188)

Ritam Jana (200798)

Utkarsh Kandi (201068)

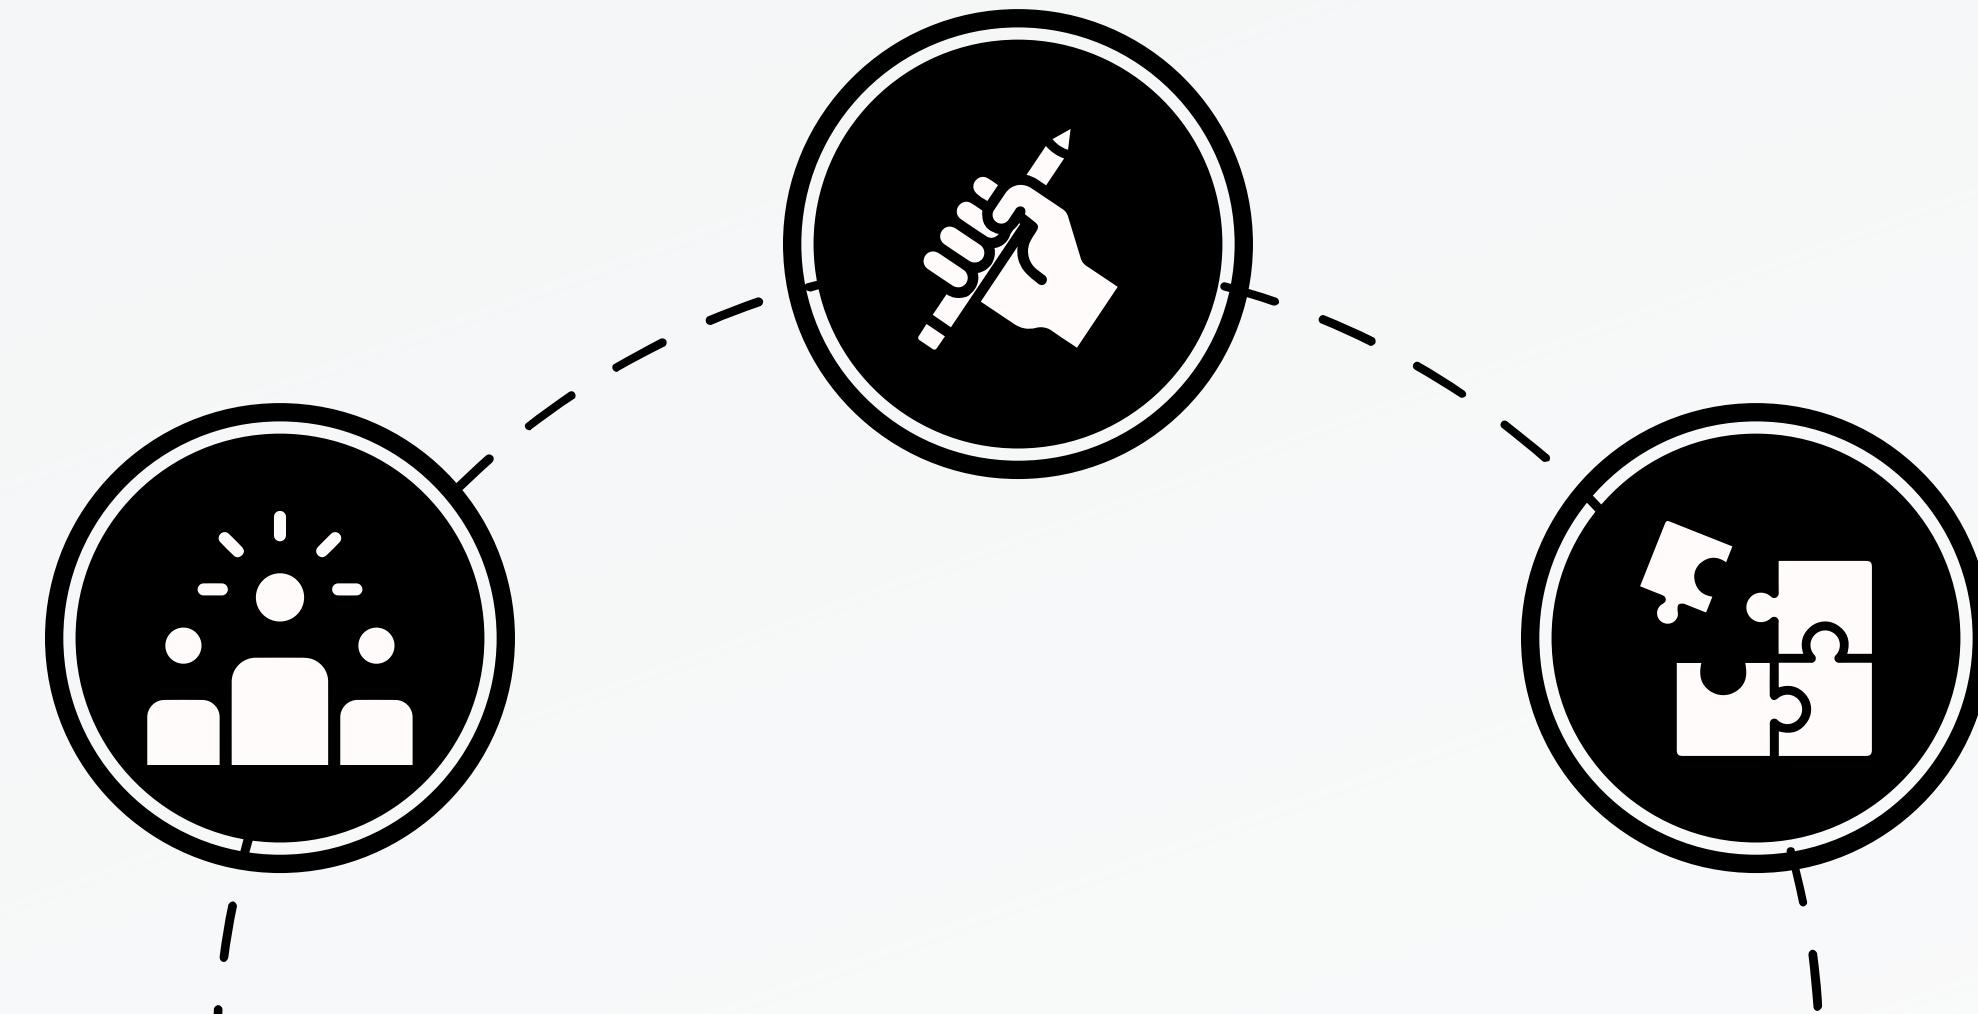
Prabuddha Singh (200691)

Aditya Palmate (200662)



GOALS AND OBJECTIVES

To introduce "Chauffeur," an open-source benchmark suite for self-driving vehicles. To highlight the motivation behind creating Chauffeur and the challenges it addresses.



LIMITATIONS OF EXISTING BENCHMARK SUITES

Traditional benchmarks:

- Conventional embedded benchmark suites (e.g., MiBench, SPEC, and PARSEC) cover various application fields, such as media processing, computer vision, animation physics, and corporate servers.
- These benchmark suites primarily focus on evaluating system/hardware design and computer architecture.

- Existing benchmarks are often specialized for either GPU workloads (e.g., Rodinia) or CPU workloads (e.g., PARSEC), resulting in a limited scope.
- This limited scope fails to accurately capture the diverse resource utilization ratios that are typical in self-driving cars.

LIMITATIONS OF EXISTING BENCHMARK SUITES

CAVBench:

- CAVBench is a benchmark suite specifically designed for assessing the performance of computing systems in the context of autonomous driving, particularly in a connected vehicle environment.
- It encompasses six key workloads relevant to self-driving cars: Simultaneous Localization and Mapping (SLAM), object detection, object tracking, battery diagnostics, speech recognition, and edge video analysis.

- CAVBench provides an in-depth analysis of the execution time breakdown for each of these applications, offering insights into the computational demands of autonomous driving tasks.
- CAVBench serves as an initial artifact to study edge computing systems for autonomous driving but fails to present a holistic view of the end-to-end self-driving scenario.

LIMITATIONS OF EXISTING BENCHMARK SUITES

Autoware:

- Autoware is a widely used open-source full-stack driving simulator built on the ROS and leverages various established open-source software libraries.
- However, Autoware and similar full-stack simulators often demand powerful hardware platforms for evaluation. For instance, for assessing Autoware includes an 8-core X86 CPU with 32GB of RAM, making it impractical for resource-constrained embedded platforms.

- To study workloads on more specialized hardware like NVIDIA DRIVE PX2 computing platform, researchers have had to make customizations to its software stack. This customization, while enabling specific evaluations, limits the ability to explore different algorithms for the same task.
- The porting of Autoware to emerging embedded platforms involves significant efforts and complexity, posing challenges for hardware designers and researchers.

LIMITATIONS OF EXISTING BENCHMARK SUITES

Apollo:

- Apollo is an industry-standard autonomous driving software framework that is developed with the primary application of self-driving vehicles in mind.
- Operating Apollo and similar frameworks on low-resource embedded hardware can be challenging, as these frameworks are designed with a focus on functionality rather than resource efficiency.

- Setting up an end-to-end driving stack for hardware and architectural studies using frameworks like Apollo is time-consuming and burdensome for those seeking to assess the hardware implications of autonomous driving systems.
- Existing benchmarks are specialized for either GPU workloads (e.g., Rodinia) or CPU. Researchers and hardware designers face difficulties in adapting such complex software.

WHY CHAUFFEUR?

Representation of Real-World Workloads

Chauffeur includes applications that accurately represent the diverse, sensor-driven, and real-time workloads encountered in self-driving systems.

End-to-End Pipeline Analysis

Chauffeur enables the analysis and optimization of the entire self-driving pipeline, allowing researchers to understand the interplay of various stages and components.

Embedded System Focus

Chauffeur addresses the unique challenges of resource-constrained embedded systems, such as real-time performance, energy efficiency, and safety.

WHY CHAUFFEUR?

Heterogeneity Support

Chauffeur accommodates the trend of utilizing GPUs and dedicated accelerators, ensuring that self-driving applications can make efficient use of available resources.

Diverse Platform Compatibility

Chauffeur is designed to work on a range of embedded platforms, reflecting the industry's use of various hardware configurations.

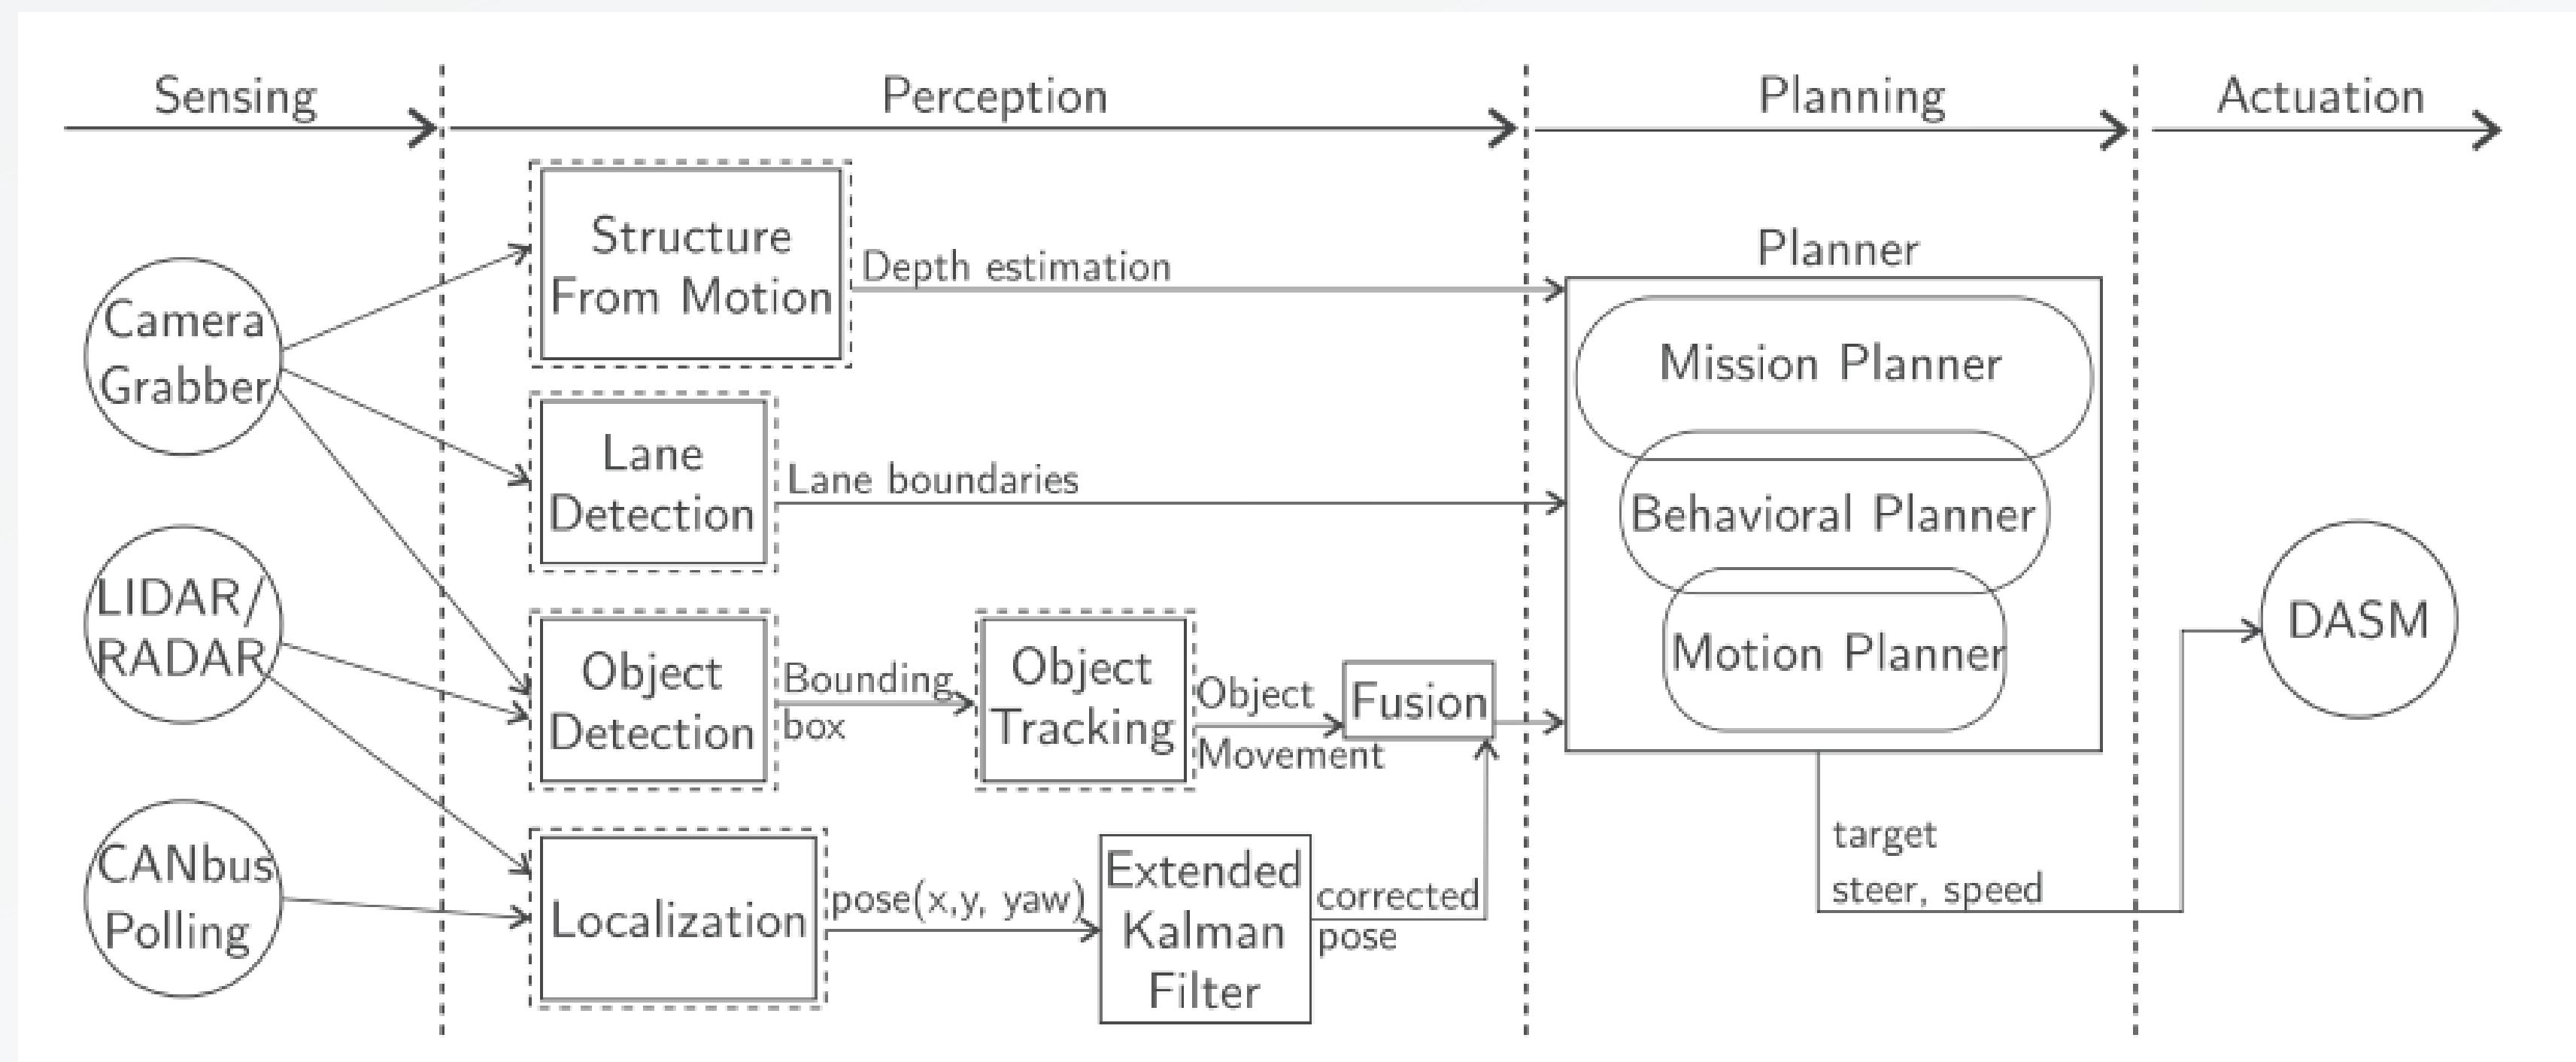
Open-Source and Research-Friendly

Chauffeur is an open-source benchmark suite that makes it easy for researchers, hardware designers, and system developers to evaluate and experiment with self-driving solutions.

COMPARISON OF POPULAR BENCHMARK SUITES

	Traditional benchmarks	CAVBench	Autoware	Apollo	Chauffeur
Self Driving Workload	✗	✓	✓	✓	✓
Configurable end to end pipeline	✗	✗	✗	✗	✓
Embedded Runtime	✓	✗	✓	✗	✓
Heterogeneity	✗	✓	✓	✓	✓
Diverse Platforms	✓	✗	✗	✗	✓
Supports research	✓	✓	✓	✓	✓

SELF-DRIVING PIPELINE



SENSORS

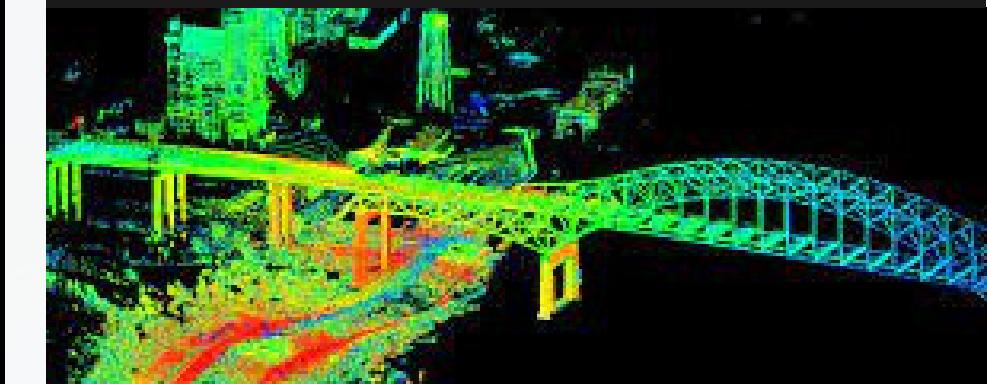
Camera



- Visual sensors are a vital component to enable the perception of the environment in self-driving vehicles.
- Require a high-bandwidth communication bus because they interface multiple cameras simultaneously.
- The camera grabber is responsible for receiving the packets from the network and storing them for later stages.

- Detection and ranging sensors are used to detect surrounding objects and calculate distances. Distance sensors assist in hazard detection and range-finding.
- These sensors are imperative during adverse weather and lighting conditions.

LIDAR/RADAR



CAN Bus Polling



- CAN bus is typically used for low-volume data transfers with high reliability.
- Used for reading Odometer values, interfacing with sensors such as IMU and controlling the steer and speed of the vehicle.

CAN BUS POLLING

- CANbus polling, also known as Controller Area Network (CAN) polling, is a communication method used to collect real-time data from hardware sensors in automotive and industrial applications
- It allows multiple devices, such as sensors, controllers, and actuators, to communicate efficiently and reliably on a shared network.

How does CANbus Polling work?

Sensor Network Setup

- Multiple sensors are connected to a CANbus network
- Each sensor is assigned a unique identifier or address (CAN ID).
- A central controller or monitoring device, such as an Engine Control Unit (ECU) in a vehicle or a Programmable Logic Controller (PLC) in an industrial setting, is also connected to the CANbus network.

Polling Process

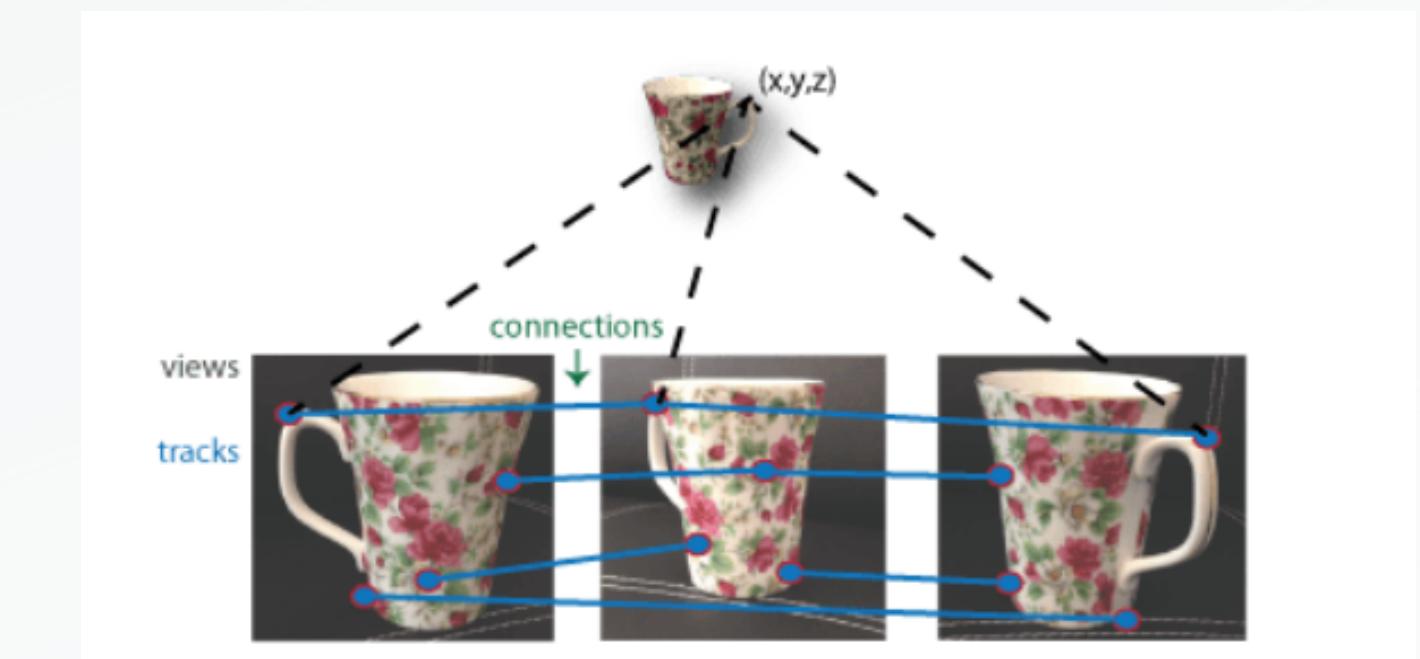
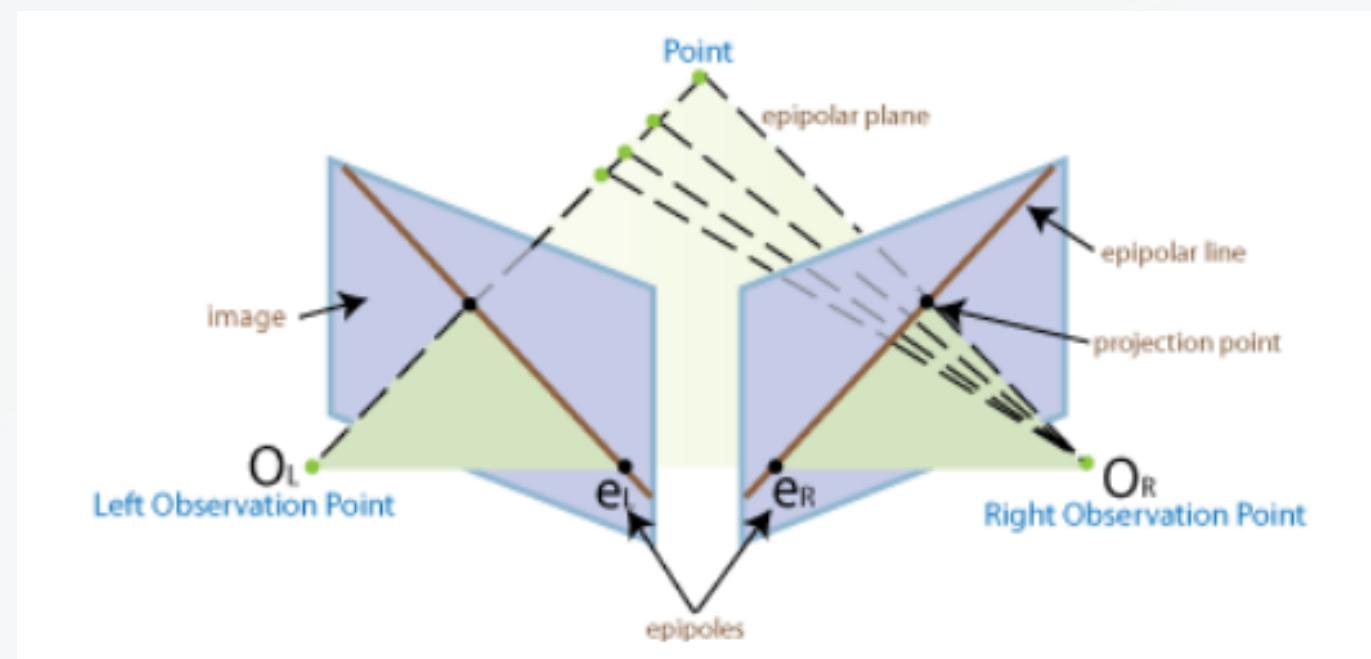
- The central controller initiates the polling process by sending a request (CAN message) with the specific CAN ID of the sensor it wants to query.
- The sensor that matches the requested CAN ID responds with its data.
- The central controller receives and processes the data.

Real Time Data collection

- The central controller can continuously or periodically poll various sensors on the CANbus network to collect real-time data.
- Each sensor responds to the controller's requests with the current sensor data, which can include information such as temperature, pressure, speed, or any other parameter it is designed to measure.
- The controller can then use this data for various purposes, including system monitoring, control, or display.

SFM - STRUCTURE FROM MOTION

- SFM is a perception application that aims to reconstruct three-dimensional structures from a sequence of two-dimensional moving images. SfM is used in many applications, such as 3-D scanning, augmented reality, and mapping (vSLAM).
- SFM uses the subsequent images to triangulate the 3D position of objects in the environment.
- SfM requires point correspondences between images. Find corresponding points by matching features or tracking points from image 1 to image 2. The fundamental matrix relates the corresponding set of points in two images from different views.



LANE DETECTION

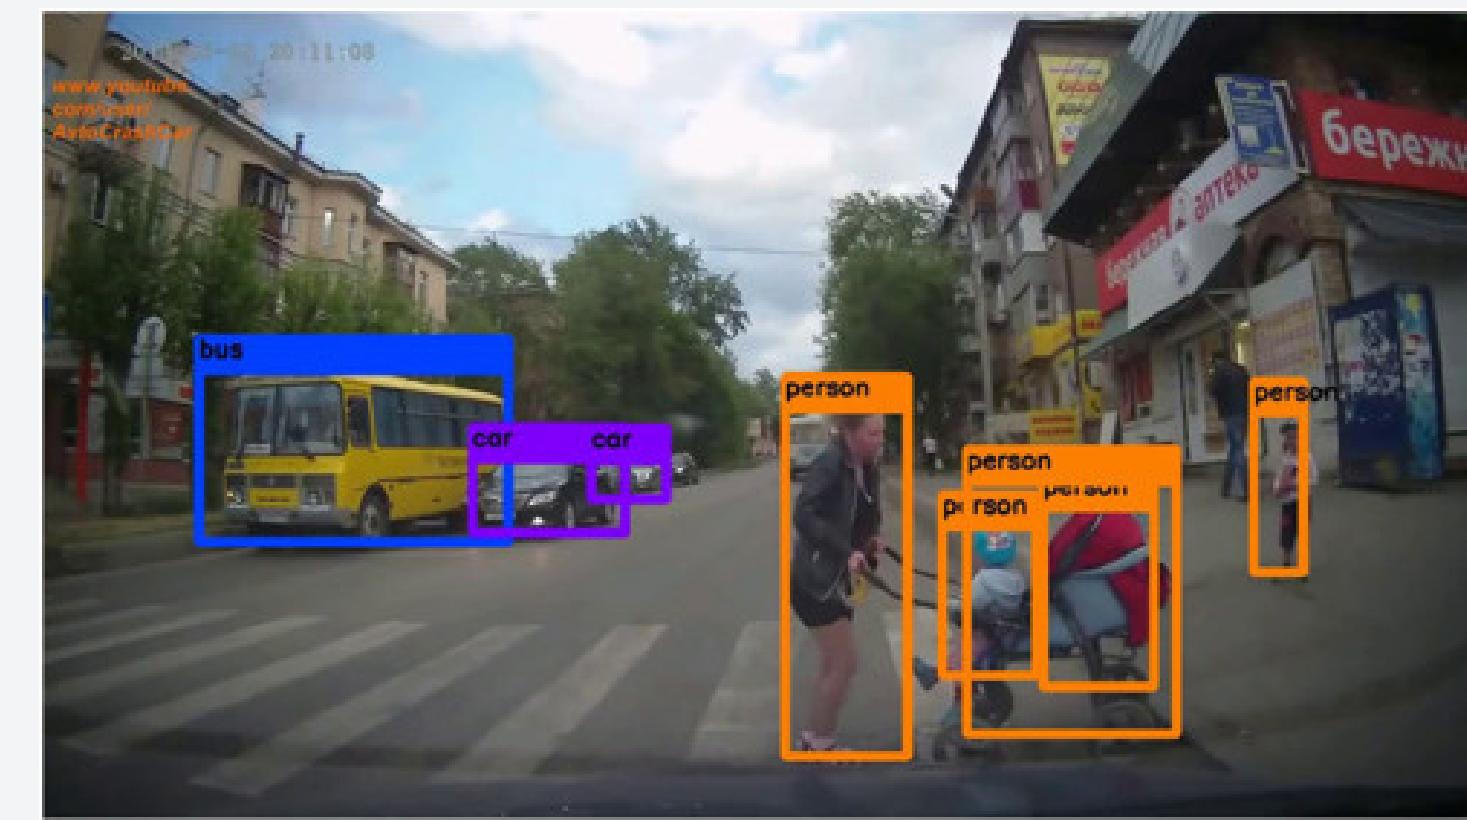
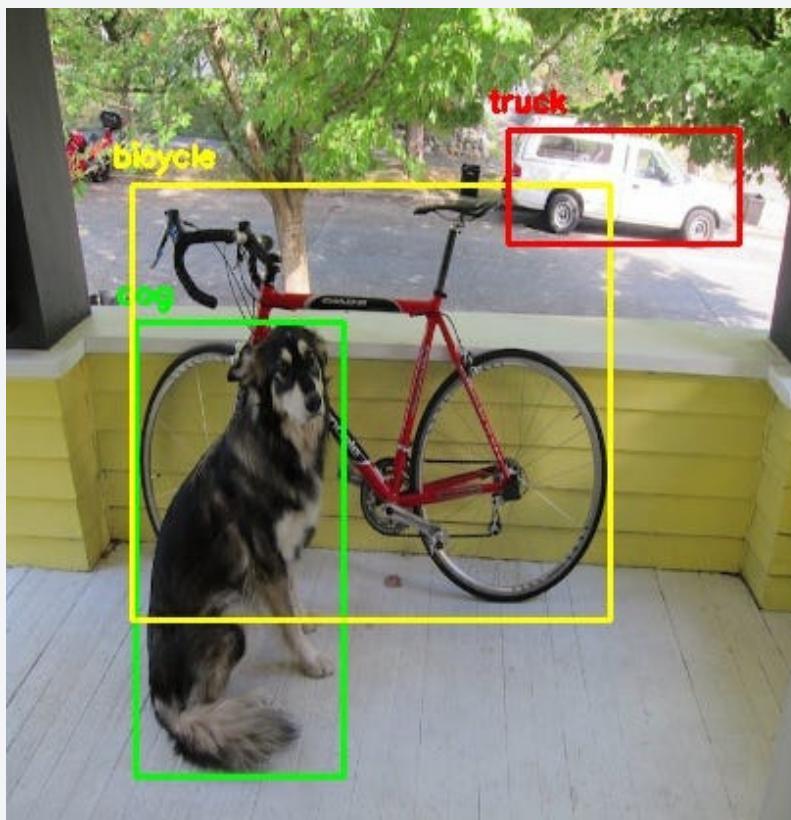
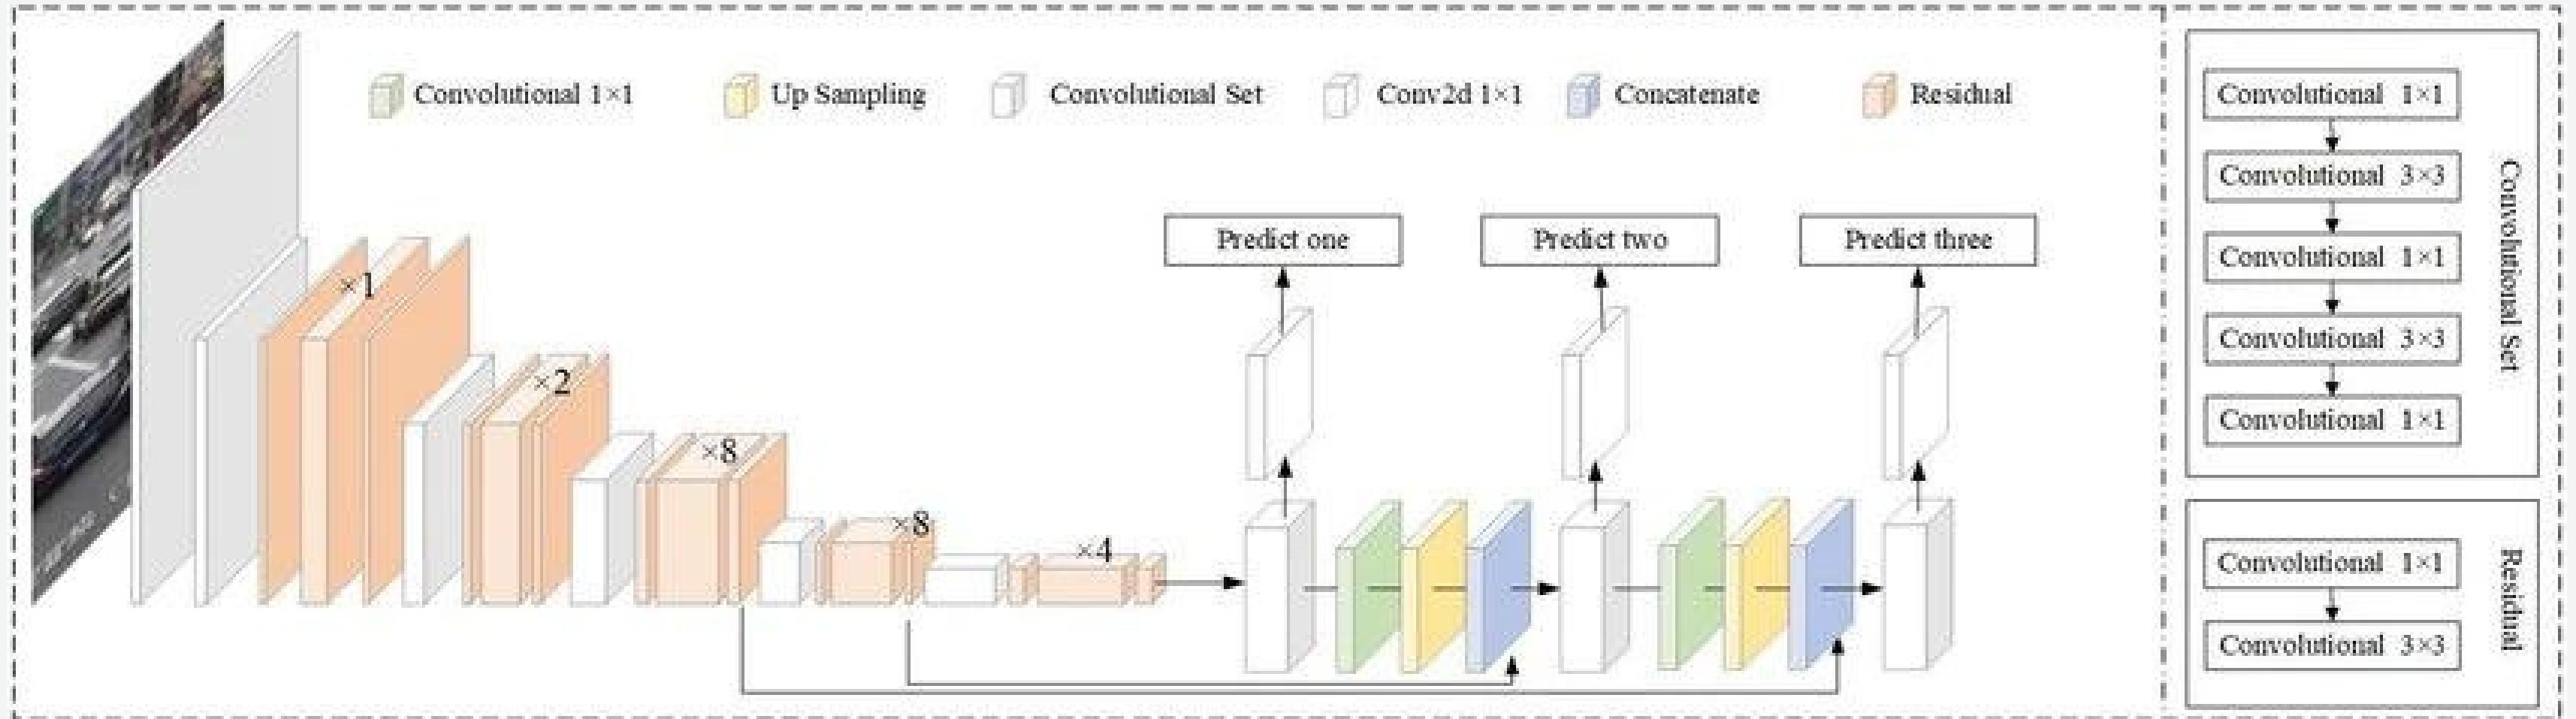
- Lane detection is a perception application that aims to detect road boundaries (lane line markings) and estimate the vehicle pose with respect to the detected lines using visual sensors on the vehicle.
- The application includes the localization of the road, the determination of the relative position between the vehicle and the road, and the analysis of the vehicle's heading direction.
- The steps followed for lane detection are:
 - Apply a distortion correction to raw images.
 - Use colour transforms to create a thresholded binary image.
 - Apply a perspective transform to generate a “bird’s-eye view” of the image.
 - Detect lane pixels and fit to find the lane boundary.
 - Determine the curvature of the lane and vehicle position with respect to the centre.
 - Warp the detected lane boundaries back onto the original image.

Applications in a Self-driving Vehicle

Application	Stage	Input	Output
Camera Grabber	➤ Sensing	➤ Packets over Automotive Ethernet	S1 ➤ Image frames in the shared memory
LIDAR/RADAR	➤ Sensing	➤ Packets over Automotive Ethernet	S2 ➤ Point cloud in the shared memory
CAN bus pooling	➤ Sensing	➤ Messages over CAN bus	S3 ➤ Sensed information in shared memory
Structure From Motion	➤ Perception	S1	P1 ➤ Depth Estimation
Lane Detection	➤ Perception	S1	P2 ➤ Lane Boundaries

OBJECT DETECTION AND TRACKING

- Vision-based object detection is a perception application that is one of the primary prerequisites for self-driving vehicles. Distance and ranging sensors (e.g., LIDAR, RADAR) alone are insufficient to meet the requirements of self-driving.
- Modern object detection applications use neural networks along with visual sensor data to identify objects in the area surrounding the vehicle by drawing bounding boxes and classifying the objects inside each bounding box.
- Given some objects of interest marked in a frame, the object tracking application locates the objects in subsequent frames in the video.
- Object tracking tracks objects as they move in the environment. It also allows the self-driving vehicle to estimate the motion of objects and predict how they will move in the subsequent frames.
- YOLOv3 is a real-time object detection algorithm that identifies specific objects in videos, live feeds, or images. The YOLO machine learning algorithm uses features learned by a deep convolutional neural network to detect an object.

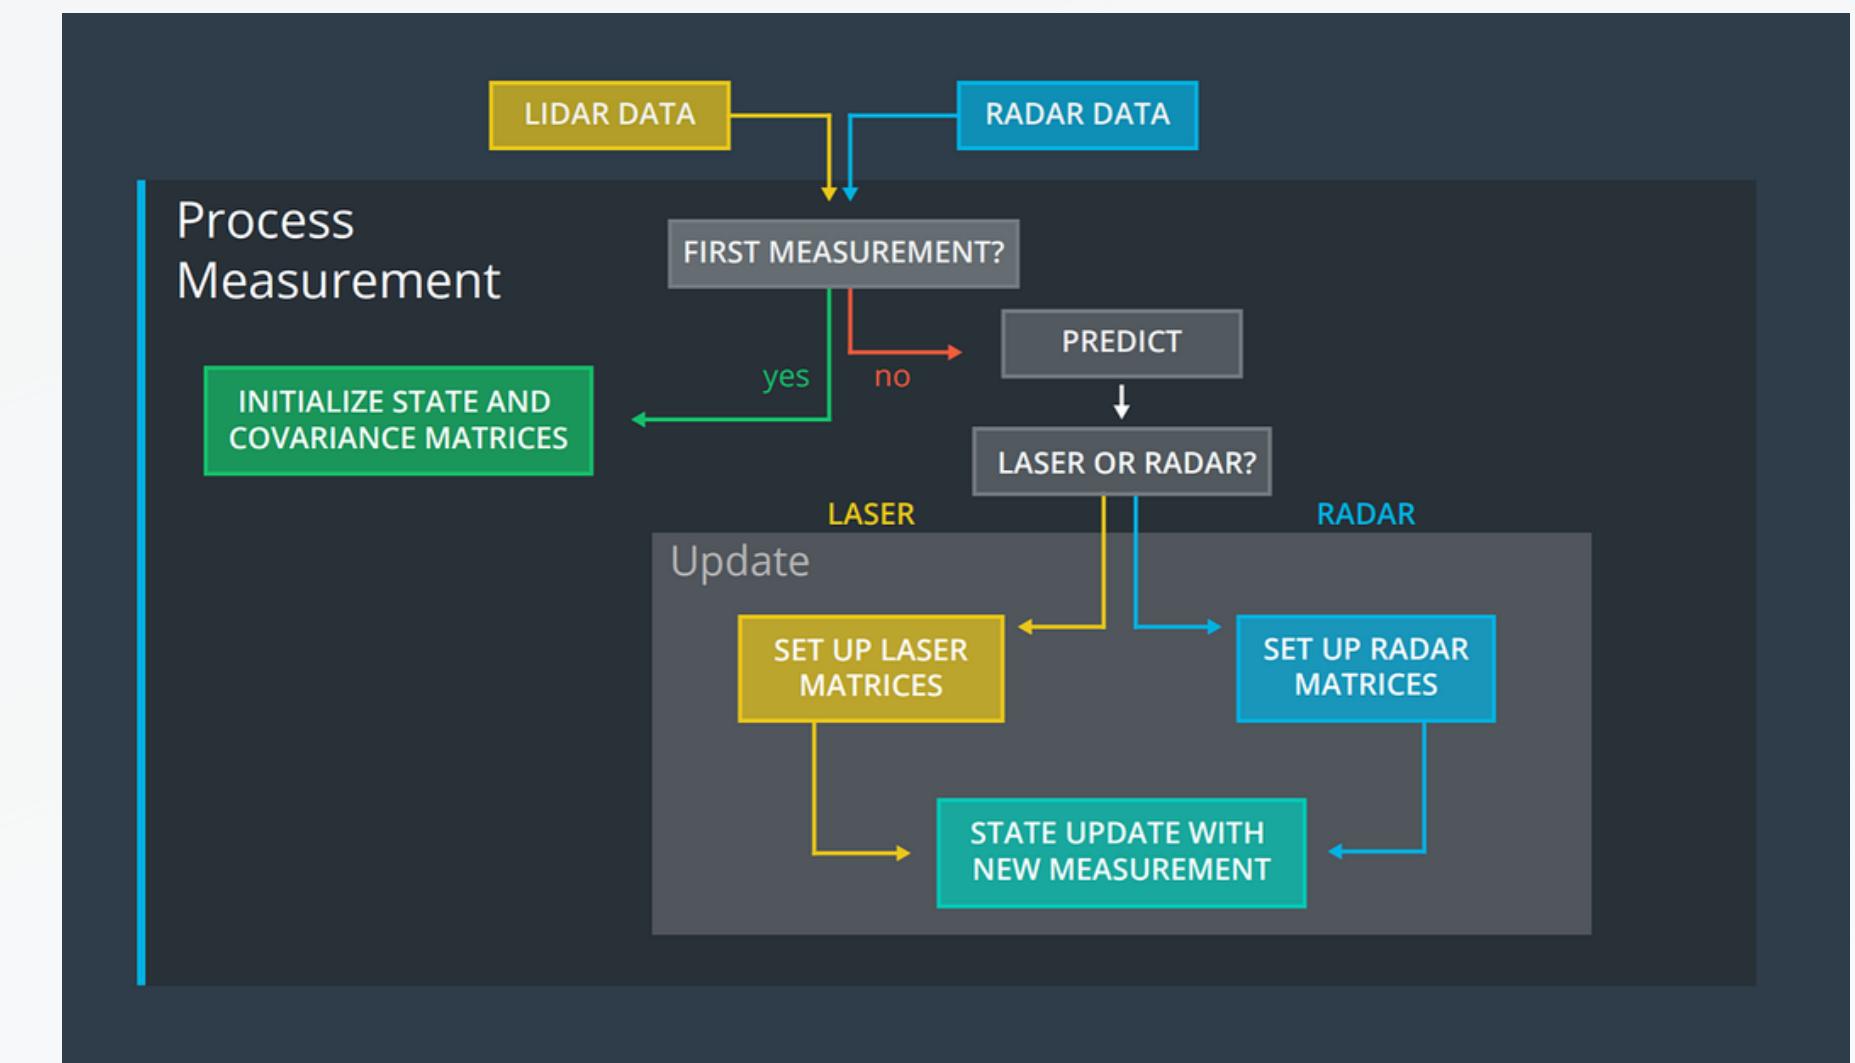


Activities Applications Tilix Mon Nov 6 00:15:34 51.1% 62.9°C 3.22 GHz 25.3% 10.29 GB 100% 1/1 + 🔍 1: /home/ritam/Documents/AUV/Software/usb/catkin_ws/src/video_stream_opencv/launch/webcam.launch http://localhost:11311 Tilix: rostopic echo /darknet_ros/bounding_boxes 2: /home/ritam/Documents/AUV/Software/usb/catkin_ws/src/darknet_ros/launch/darknet.launch http://localhost:11311 3: rostopic echo /darknet_ros/bounding_boxes

```
[DEBG] [1699197576.390765263, 15.916000000]: DiffDrive(ns = //): <odometryTopic> = odom
[FPS:1.3] [1699197576.390779097, 15.916000000]: DiffDrive(ns = //): <odometryFrame> = odom
[Objects: [1699197576.390796771, 15.916000000]: DiffDrive(ns = //): <robotBaseFrame> = base_link
[DEBUG] [1699197576.390857471, 15.916000000]: DiffDrive(ns = //): <publishWheelTF> = true
[cup: 81% [1699197576.390873109, 15.916000000]: DiffDrive(ns = //): missing <publishOdomTF> defau
[book: 33%]
[DEBUG] [1699197576.390893682, 15.916000000]: DiffDrive(ns = //): <publishWheelJointState> = tr
[DEBUG] [1699197576.390948272, 15.916000000]: DiffDrive(ns = //): <wheelSeparation> = 0.2999999
[DEBUG] [1699197576.390963298, 15.916000000]: DiffDrive(ns = //): <wheelDiameter> = 0.080000000
[DEBUG] [1699197576.390977157, 15.916000000]: DiffDrive(ns = //): <wheelAcceleration> = 1.8
[DEBUG] [1699197576.390990761, 15.916000000]: DiffDrive(ns = //): <wheelTorque> = 30
[DEBUG] [1699197576.391004755, 15.916000000]: DiffDrive(ns = //): <updateRate> = 100
[WARN] [1699197576.391049205, 15.916000000]: DiffDrive(ns = //): missing <odometrySource> defa
ult is 1
[DEBUG] [1699197576.391078705, 15.916000000]: DiffDrive(ns = //): <leftJoint> = front_left_whee
l_joint
[DEBUG] [1699197576.391092971, 15.916000000]: DiffDrive(ns = //): <rightJoint> = front_right_wh
eel_joint
[INFO] [1699197576.391093692, 15.916000000]: DiffDrive(ns = //): Advertise joint states
[INFO] [1699197576.391093692, 15.916000000]: nsecs: 689224105393441297, 15.916000000]: Subscri
be to cmd_vel
[frame_id: "webcam_optical_frame" 916000000]: DiffDrive(ns = //): Advertise odom on odom
[bounding_boxes: [1699197576.404735668, 15.927000000]: Trying to publish message of type [nav_msgs/Odome
tr-/cd5e73d190d741a2f92e81eda573aca7] on a publisher with type [nav_msgs/Odometry/cd5e73d190d74
1a2f probability: 0.7084612488746643
[DEBUG] [1699197576.404866562, 15.927000000]: Trying to publish message of type [sensor_msgs/Jo
intS ymin: 125dc d76a6cfaef579bd0f34173e9fd] on a publisher with type [sensor_msgs/JointState/306
5dc xmax: 278579bd0f34173e9fd]
[ymin: 125dc ymax: 274] killing on exit
[gaz id: 41 killing on exit
[DEBUG] [1699197576.5952599719641, 66.252000000]: Unloaded
[DEBUG] [1699197652.599719641, 66.252000000]: Unloaded
[gaz probability: 0.3303779363632202
[ros xmin: 330 killing on exit
[mas ymin: 135 ing on exit
[shut xmax: 640 processing monitor...
... ymax: 341 down processing monitor complete
done id: 73
    Class: "book"
--> ~/Doc/A/Software/Simulation/assignment1/catkin_ws/src > git master !11 74 1m 51s
```

EXTENDED KALMAN FILTER (EKF)

- The result of localization (using distance sensors like RADAR) is prone to drift over time and can be noisy.
- A Kalman filter is an excellent candidate for combining distance information with other vehicle status information (e.g., IMU, Odometer, GPS) and handling such disturbances.
- Kalman filter fuses multiple data sources and performs continuous prediction (for missing data) and correction (for drifting data) on the localization results.
- The extended Kalman filter (EKF) is the nonlinear version of the Kalman filter.



FUSION

- The fusion task helps combine object information with the car's location on the fly. The information is sourced by pre-processing raw data from different sensors (e.g., cameras, different types of RADAR, LIDAR) and fused synchronously.
- The fusion process involves transforming different coordinate systems and updating the environment maps periodically in real-time.



PLANNING

Path Planner

- The planning stage is responsible for understanding higher-level goals from the user and converting them to purposeful decisions to achieve the higher-level goals while avoiding obstacles.
- The complexity of this stage compels a hierarchical design by partitioning the software into layers:
 - Mission planning: Mission planning computes the global trajectory based on the current location and the target destination
 - Behavioural planning: Behavioural planning generates local objectives
 - Motion planning: Motion planning takes the local objectives and generates the control plan to actuate the steering and speed.

DASM

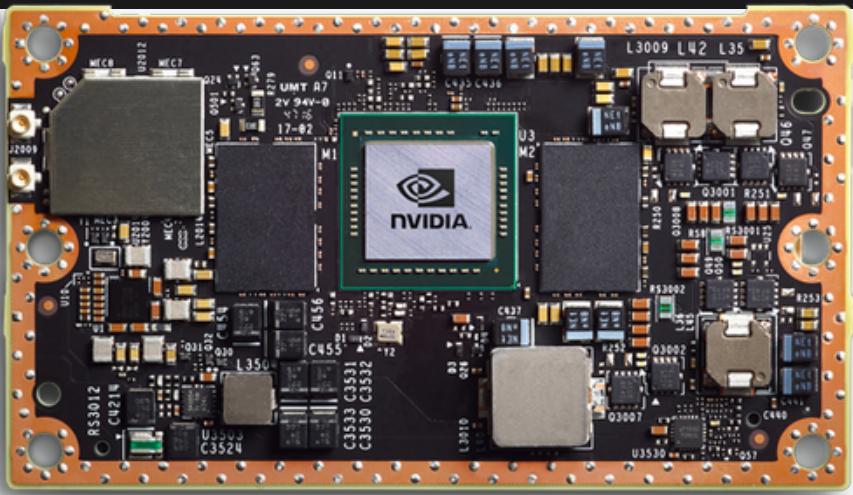
- The **Driver Assistance System Module** performs the final actuation stage in the pipeline. The local objectives of the path planner are realized through a PID controller.
- The PID controller actuates the speed and steering based on the velocity and angle commands and can automatically use the odometer and IMU feedback to maintain the targets set by the path planning stage.

Applications in a Self-driving Vehicle

Application	Stage	Input	Output
Object Detection	➤ Perception	S1 S2	I1 ➤ Bounding Box
Object Tracking	➤ Perception	I1	I2 ➤ Object Movement
Localization	➤ Perception	S2	I3 ➤ Position and orientation pose
Extended Kalman Filter	➤ Perception	I3 S3	I4 ➤ Corrected pose
Fusion	➤ Perception	I2 I4	P3 ➤ Fused object and vehicle location
Path Planner	➤ Planning	P1 P2 P3	A1 ➤ Spatio-temporal trajectory
Driver Assistance System Module	➤ Actuation	A1	➤ Steer, Brake

SUPPORTED PLATFORMS

JETSON TX2 GPU



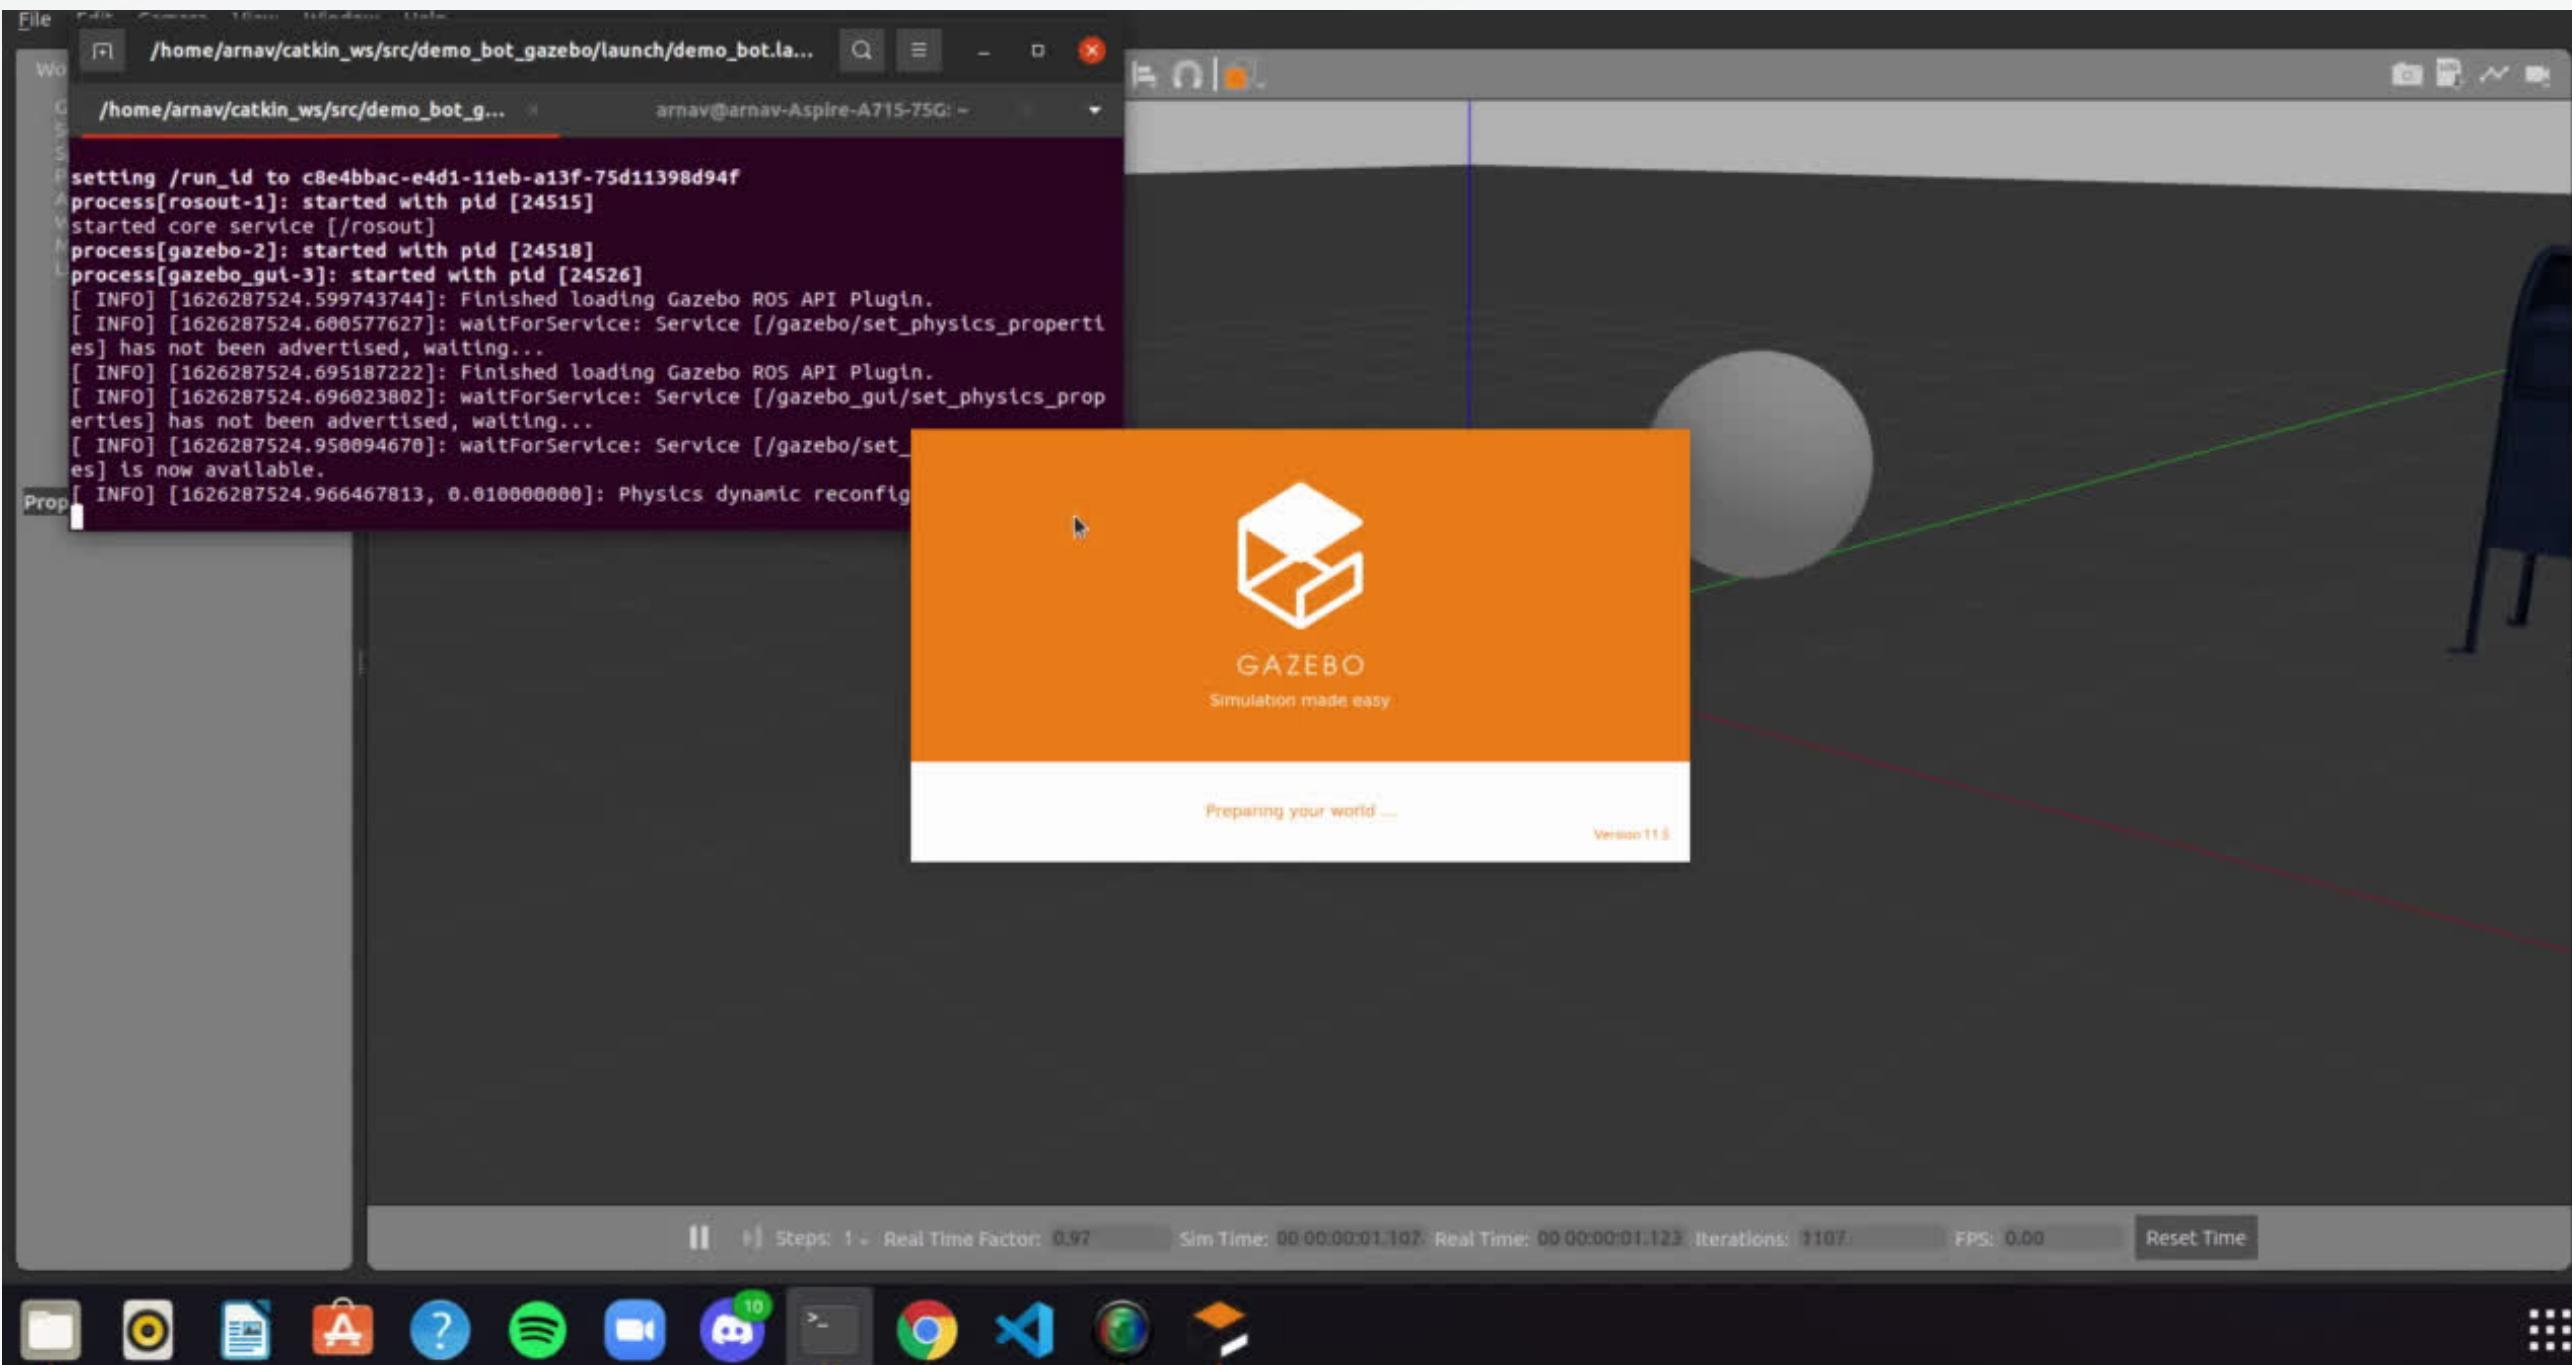
- 8 GB 128-bit LPDDR4 Memory
- 7.5 Watt
- 59.7 GB/s of memory bandwidth
- 1.33 TFLOPS

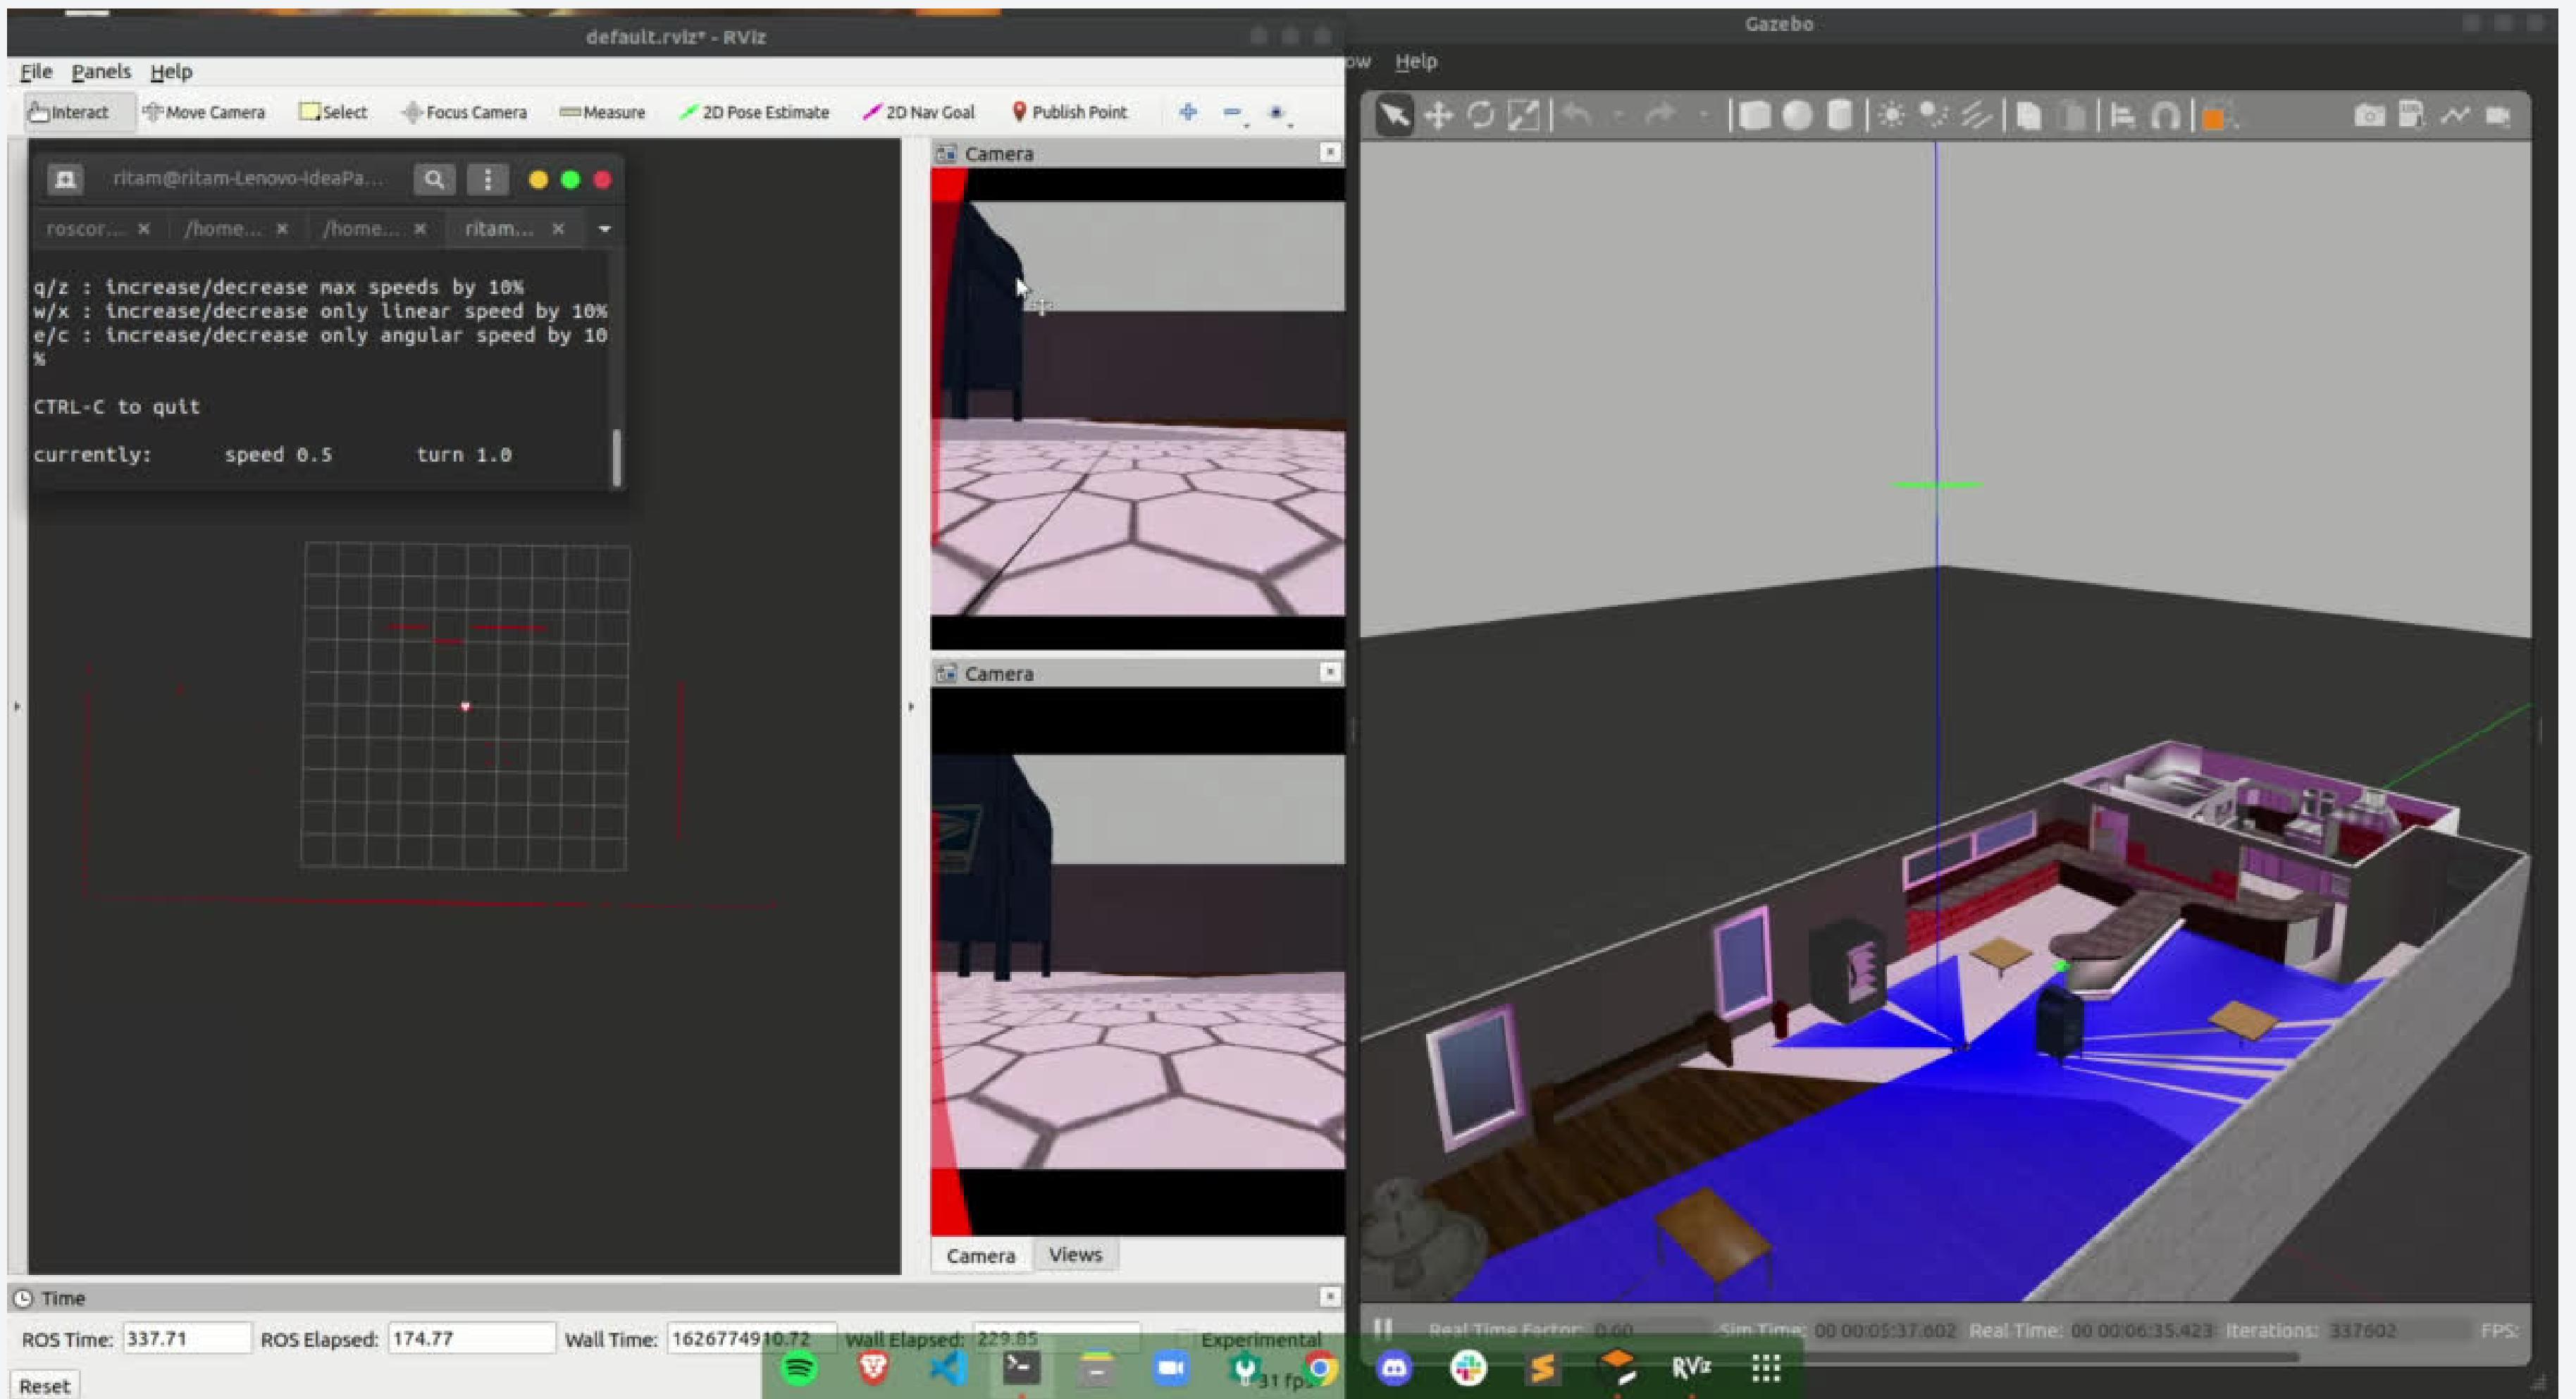
NVIDIA Drive PX2

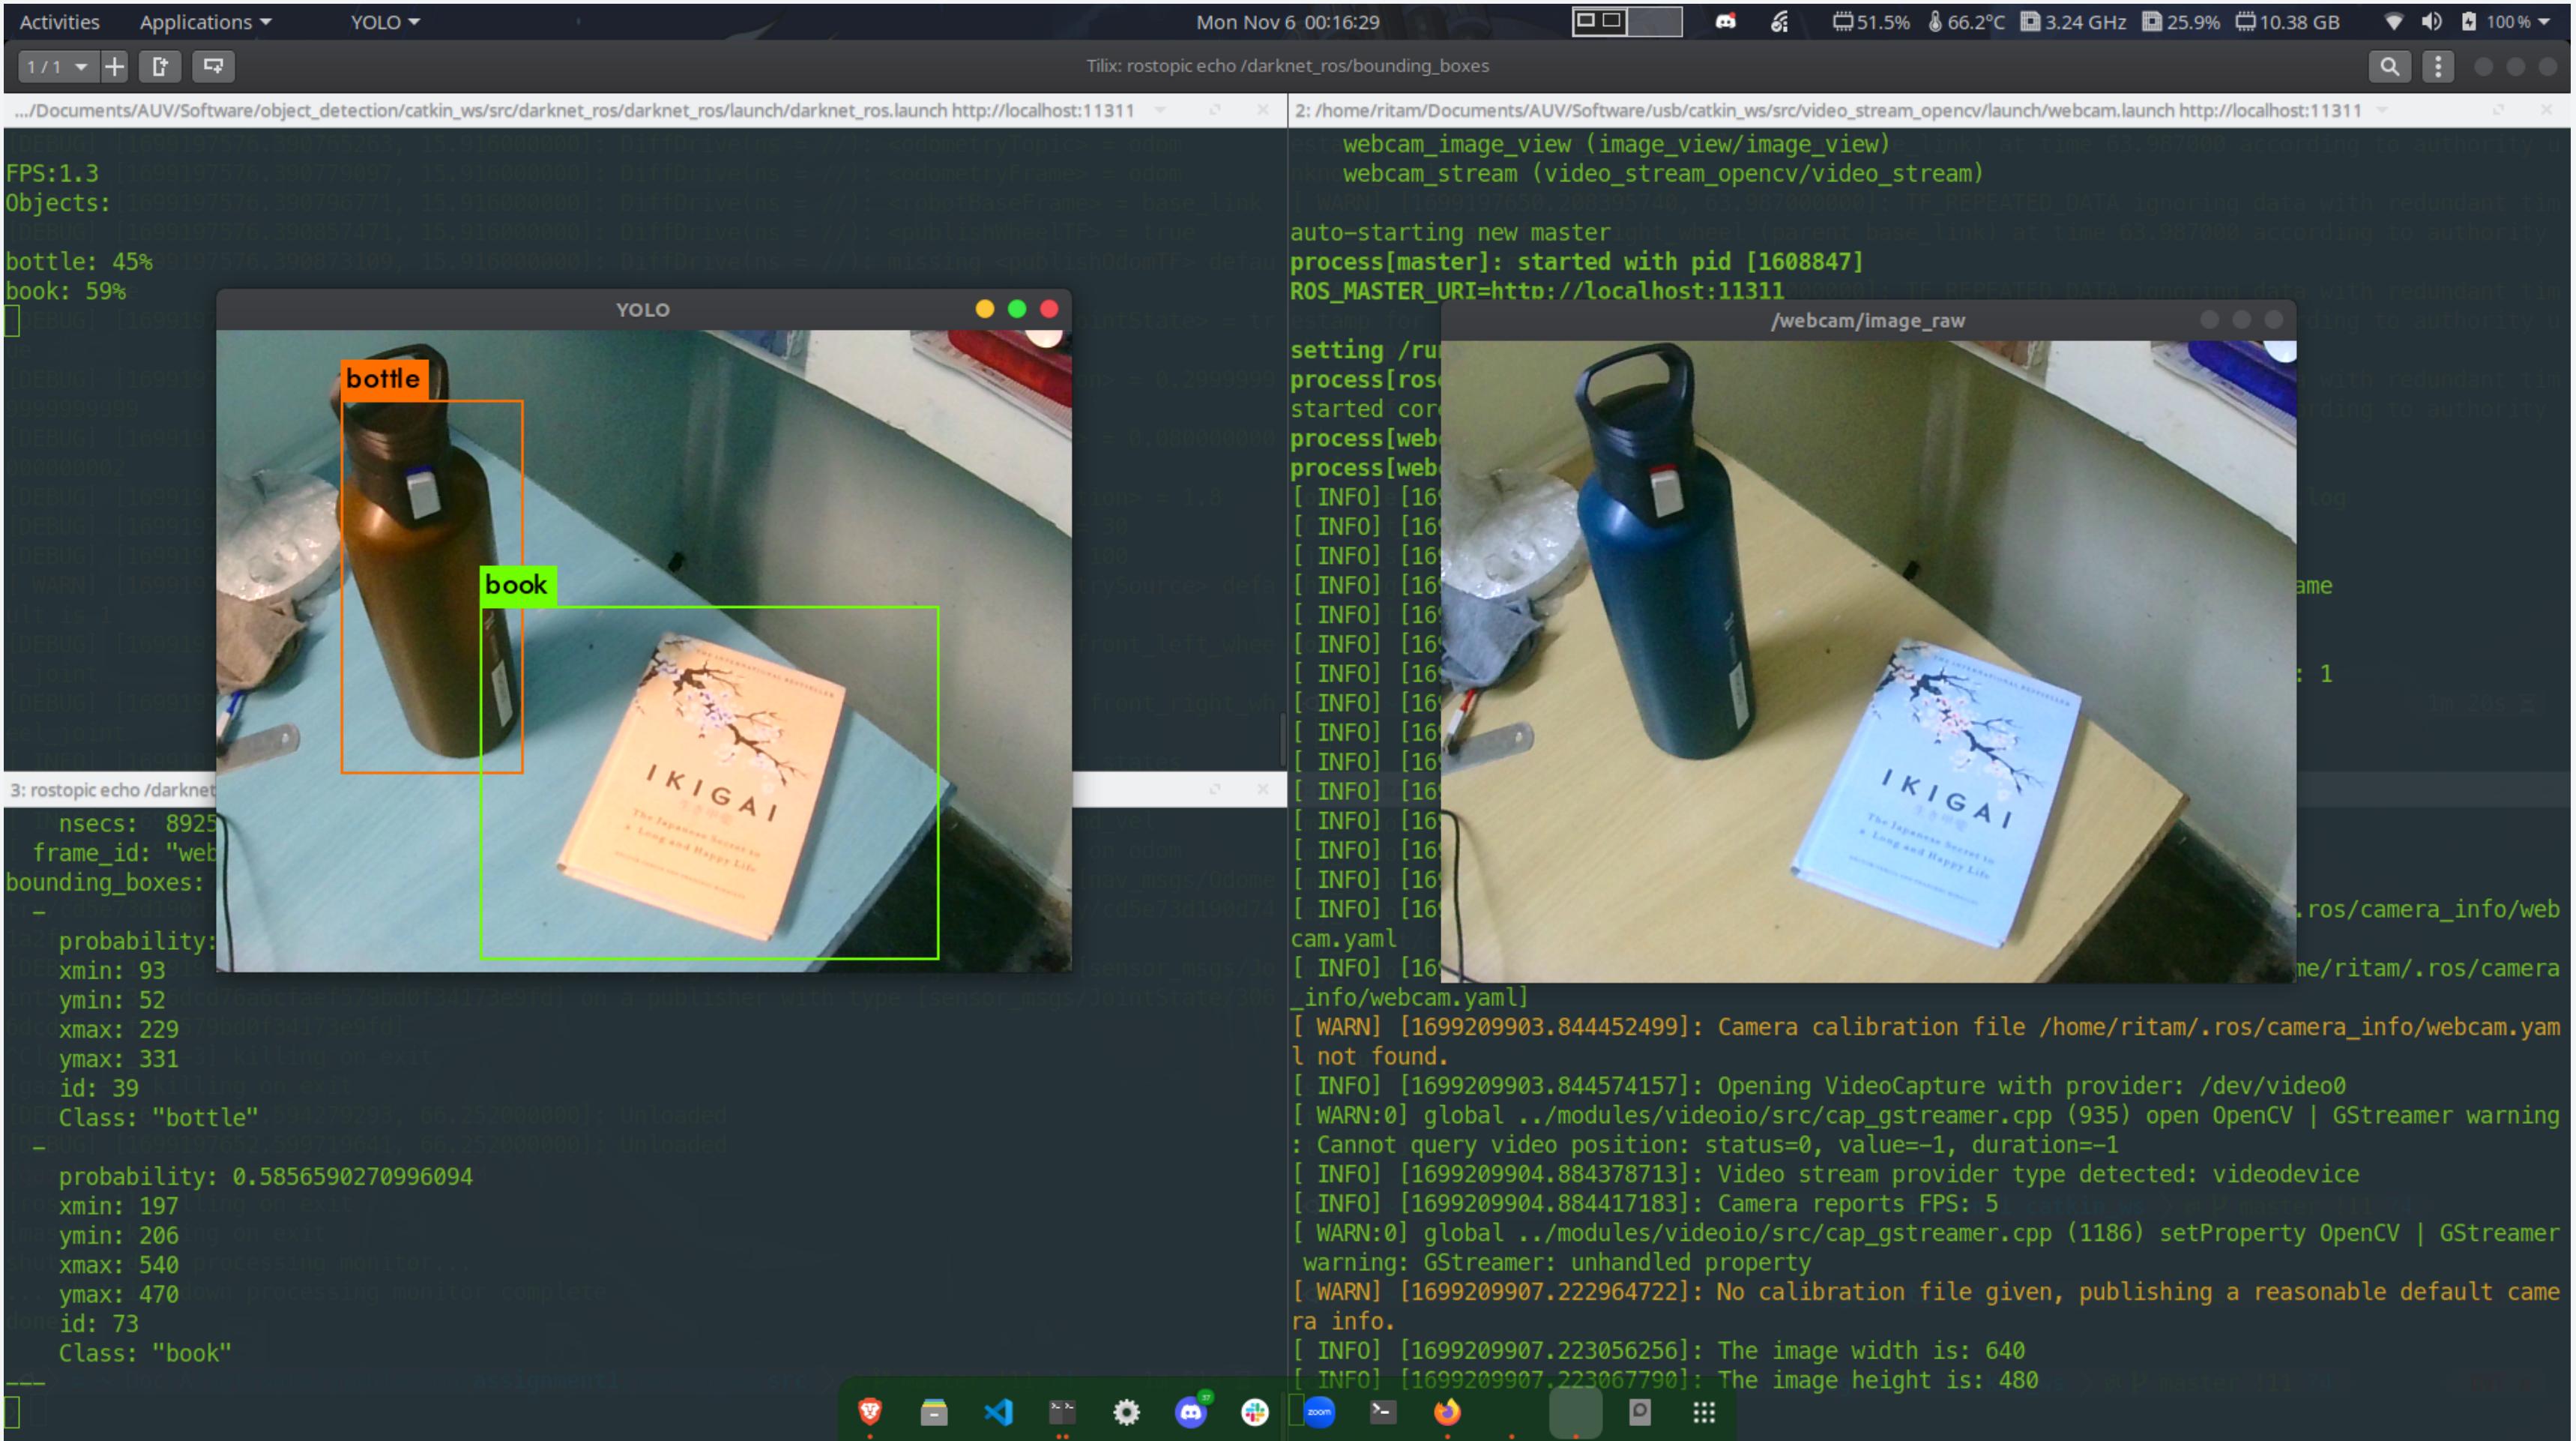
- No cross-compiler support
- 250 Watt Total Dissipated Power (TDP)
- 12 core ARM CPU (in total)



IMPLEMENTATION







Definition of the SLAM Problem

Given

- The robot's controls

$$u_{1:T} = \{u_1, u_2, u_3, \dots, u_T\}$$

- Observations

$$z_{1:T} = \{z_1, z_2, z_3, \dots, z_T\}$$

Wanted

- Map of the environment

$$m$$

- Path of the robot

$$x_{0:T} = \{x_0, x_1, x_2, \dots, x_T\}$$

Kalman Filter Algorithm

```
1: Kalman_filter( $\mu_{t-1}, \Sigma_{t-1}, u_t, z_t$ ):  
2:    $\bar{\mu}_t = A_t \mu_{t-1} + B_t u_t$   
3:    $\bar{\Sigma}_t = A_t \Sigma_{t-1} A_t^T + R_t$   
4:    $K_t = \bar{\Sigma}_t C_t^T (C_t \bar{\Sigma}_t C_t^T + Q_t)^{-1}$   
5:    $\mu_t = \bar{\mu}_t + K_t(z_t - C_t \bar{\mu}_t)$   
6:    $\Sigma_t = (I - K_t C_t) \bar{\Sigma}_t$   
7:   return  $\mu_t, \Sigma_t$ 
```

SUMMARY



We implemented a simulation of a bot in a gezebo world with cameras, LIDAR to map and track the environment.



We implemented a YOLOv3 based object detection and tracking which was able to identify objects with an accuracy of 91%



We also implemented a Kalman filter-based SLAM algorithm to map a path avoiding obstacles from source to destination.

**THANK
YOU**

